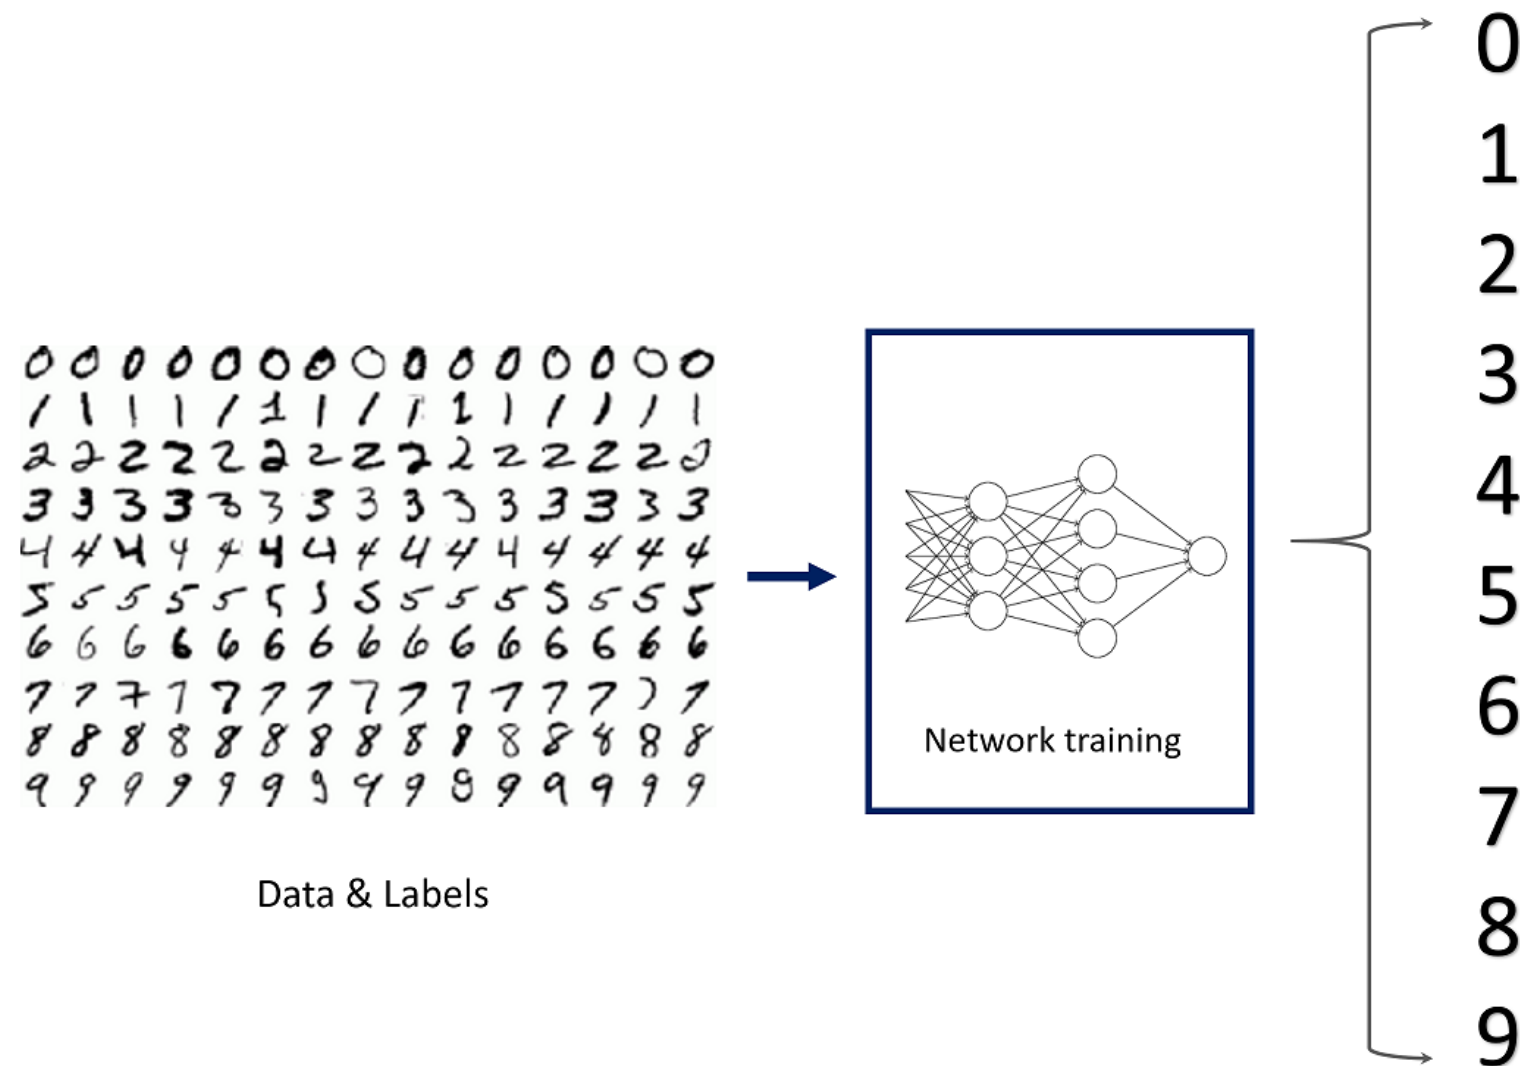


Classifying MNIST digits using MLP

TA. Hwanmoo Yong



Modified National Institute of Standards and Technology database

- 60,000 images for the training and 10,000 images for the test.
- Size : 28x28 pixels
- Only 1 channel (Gray scale)

Preparing Data

1. Load

- Use easy loading function in Keras library
- Need to create custom loading function when trying to load your own dataset.

2. Reshape

- Change size/shape of the dataset array

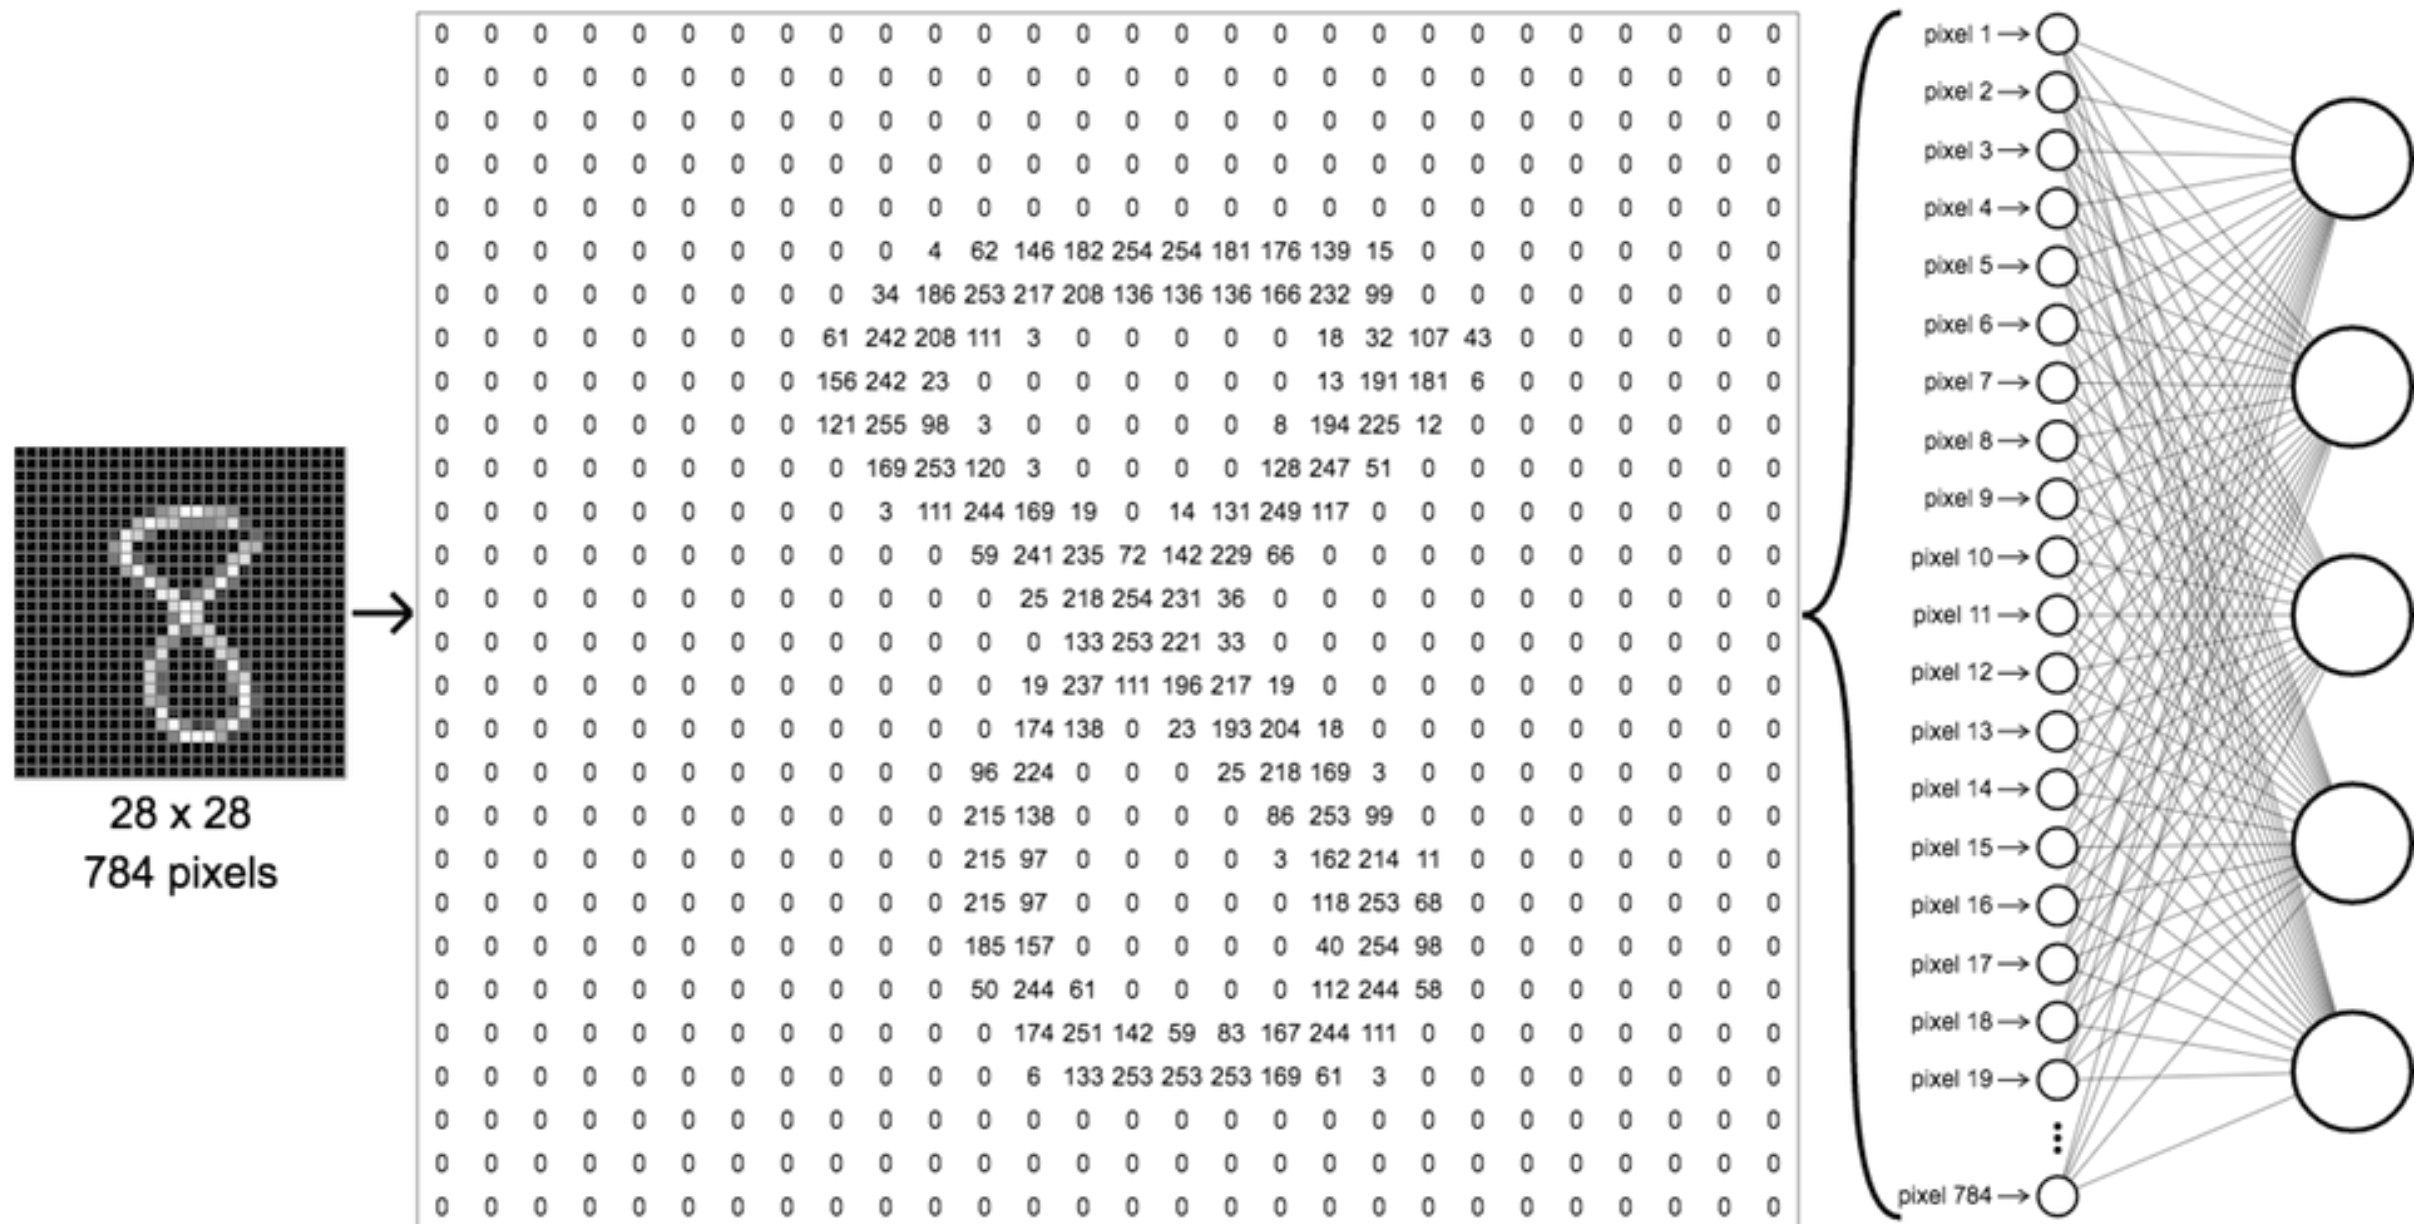
3. Normalize

- Standardize features by removing the mean and scaling to unit variance
- Divide by the maximum value

4. Convert to one-hot encoding/vector

- Use `keras.utils.to_categorical` function
- Need when training a neural network which classifies given data

Reshape




- Reshape 28x28 2-d vector to 1-d vector (size 764)

```
13 x_train = x_train.reshape(60000, 784)
14 x_test = x_test.reshape(10000, 784)
```

Normalize

- required when features have different ranges.





Common data preprocessing pipeline. Left: Original toy, 2-dimensional input data. Middle: The data is zero-centered by subtracting the mean in each dimension. The data cloud is now centered around the origin. Right: Each dimension is additionally **scaled by its standard deviation**. The red lines indicate the extent of the data - they are of unequal length in the middle, but of equal length on the right.

Normalize

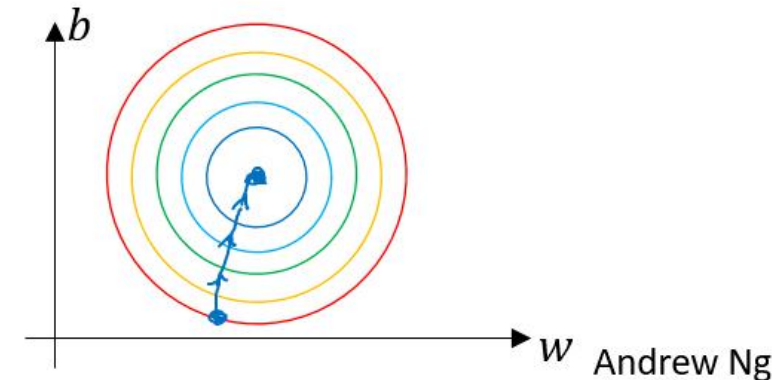

Why normalize inputs?

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Unnormalized:
 $w_1: x_1: 1 \dots 1000 \leftarrow$
 $w_2: x_2: 0 \dots 1 \leftarrow$
 $-1 \dots 1$



Normalized:



Andrew Ng

```
15 x_train = x_train.astype('float32')
16 x_test = x_test.astype('float32')
17 x_train /= 255
18 x_test /= 255
```


#Select numerical columns which needs to be normalized

```
train_norm = x_train[x_train.columns[0:10]]
test_norm = x_test[x_test.columns[0:10]]
```

Normalize Training Data

```
std_scale = preprocessing.StandardScaler().fit(train_norm)
x_train_norm = std_scale.transform(train_norm)
```


Convert to one-hot encoding/vector



```
22 # convert class vectors to binary class matrices
23 y_train = keras.utils.to_categorical(y_train, num_classes)
24 y_test = keras.utils.to_categorical(y_test, num_classes)
```

Creating a model

1. Sequential (keras)

- Easy way of creating a network.
- Convenient, but hard to use when creating complex networks.

2. Dense Layer

- Fully connected layer.
- Consists of weights and biases

3. Activation Functions

- Defines the output of the node
- ReLU, tanh, sigmoid



4. Dropout Layer

- Prevents overfitting
- Randomly removes the connections between weights

4. Softmax

- Converts the scores to the probabilities
- Usually locates at the end of the network

Dense Layer



$$y_i = x_j W_{i,j} + b_i$$


Where:

- x_j - j_{th} value in input tensor.
- y_i - i_{th} value in output tensor.
- W_{ij} - weight of j_{th} input element for i_{th} neuron.
- b_j - bias for i_{th} neuron.


```

29 model = Sequential()
30 model.add(Dense(512, activation='relu', input_shape=(784,)))
31 model.add(Dropout(0.2))
32 model.add(Dense(512, activation='relu'))
33 model.add(Dropout(0.2))
34 model.add(Dense(num_classes, activation='softmax'))
35
  
```


Activation Functions



Sigmoid





Tanh




ReLU (Rectified Linear Unit)

- Defines the output strength of the signal to the following layer
- Select proper and better activation functions by trial and errors

Overfitting




(a) Standard Neural Net



(b) After applying dropout.

- Prevents overfitting when the network has many hidden layers
- Skips / removes the connections between weights

Softmax



$$P(y=j \mid \theta^{(i)}) = \frac{e^{\theta^{(i)}}}{\sum_{j=0}^k e^{\theta_k^{(i)}}}$$

Softmax function

where $\theta = w_0x_0 + w_1x_1 + \dots + w_kx_k = \sum_{i=0}^k w_ix_i = w^T x$

- Converts the scores to the probabilities
- We select the index of the largest probability as a classification result

Training a network

1. Compile

- Sets loss/optimizer
- Use ‘categorical_crossentropy’ in classification project.


2. Fit

- Actual training part.
- Need to put train, validation data and other parameters for the training

3. Evaluate

- Gives the loss and the accuracy of the classification results with given dataset
- Use ‘predict’ instead when the actual classification result is needed.

CROSS-ENTROPY



- Calculate how the predictions of the network are different from the real values
- ie. Cross entropy, Mean squared error, Mean absolute error

Training a network

```
38 model.compile(loss='categorical_crossentropy',  
39               optimizer=RMSprop(),  
40 ▼           metrics=['accuracy'])  
41  
42 history = model.fit(x_train, y_train,  
43                   batch_size=batch_size,  
44                   epochs=epochs,  
45                   verbose=1,  
46                   validation_data=(x_test, y_test))  
47 score = model.evaluate(x_test, y_test, verbose=0)  
48 ▼ print('Test loss:', score[0])  
49 ▼ print('Test accuracy:', score[1])
```